

# Deploying ZKP Frameworks at Scale: Challenges and Proposed Solutions

Piergiuseppe Mallozzi  
mallozzi@berkeley.edu  
Berkeley, USA

## Abstract—

Zero-knowledge proofs (ZKPs) are a powerful tool for privacy-preserving computation, but their deployment is limited due to the lack of trust in how sensitive data is handled, difficulties in constructing proofs, and scalability concerns. In this paper, we propose a blockchain-integrated approach, which addresses the critical challenges in ZKP-based applications. The proposed solution focuses on establishing trust in data handling through verifiable data provenance and access control policies while keeping zero-knowledge on sensitive data. We also tackle the process of circuit generation by providing a user-friendly interface for compiling circuits at scales. Finally, we evaluate the scalability of our approach through a set of experiments on growing data structure and circuit complexity.

## I. INTRODUCTION

Zero-knowledge proofs (ZKPs) allow one party to prove knowledge of a statement to another party without revealing any additional information beyond the statement's validity. ZKPs have a wide range of applications, including authentication, identity management, and data privacy. However, the current deployment of ZKP frameworks is limited due to the lack of trust in how data is handled, the difficulty in constructing the proofs, and scalability concerns.

In this paper, we propose solutions to address these challenges. We tackle three critical issues in ZKP-based applications. First, we address the problem of establishing trust in how the required data for producing and verifying the proofs is handled. In particular, handling sensitive data requires ensuring transparent and verifiable authenticity and provenance of data, implementing access controls based on data provider policies, and maintaining a public log for transparency. We propose a general statement verification process that includes the proof of data provenance by design and an architecture of smart contracts to handle access control policies and the on-chain verification of proofs.

Second, we address the problem of circuit generation, which involves creating an arithmetic circuit to formulate the desired function to be applied to the data. Function generation is a challenging task that requires transparency and agreement on the process to establish trust and accountability. To simplify this process, our tool provides a user-friendly interface for generating circuits at scale. Our tool accepts a variety of data formats and a library of functions that can be chosen to analyse the data. The tool ultimately compiles down to circuits in the Noir language [1], a newly released Zero-Knowledge framework by Aztec.

Lastly, we evaluate the scalability of our approach and address how our approach handles data of growing size and circuits of growing complexity.

We have implemented our approach in a platform named FACT FORTRESS, a blockchain-integrated solution that utilizes ZKPs to enable efficient and secure *fact-checking* with transparent and verifiable authenticity of data without compromising privacy. The platform is available open-source, and it includes a front-end and a separate tool for circuit generation (code available<sup>1</sup>)

## II. BACKGROUND

*a) Zero-knowledge proof:* is a cryptographic technique that allows one party (the prover) to prove to another party (the verifier) that they know a particular piece of information, without revealing that information itself. This can be especially useful in scenarios where sensitive data needs to be kept confidential to ensure the *privacy* and *security* of the data.

The process of zero-knowledge proof involves a prover and a verifier, who must agree on the function to be executed on the data. Let  $f(w, x) = y$  be a function on the inputs  $w$  and  $x$  where  $w$  (often called the *witness*) is private and  $x$  and  $y$  are public. The prover must produce a proof  $\pi$  that convinces the verifier that they know a secret input  $w$  such that  $f(w, x) = y$ .

Figure 1 illustrates the interaction between the prover and verifier to prove a generic claim  $f(w, x) = y$  without revealing any information about  $w$ . The prover takes as input private data  $w$  and public data  $x$  and  $y$ , and generates a proof  $\pi$ . The verifier takes as input public data  $x$ ,  $y$ , and  $\pi$ , and can either accept or reject the proof based on whether it is valid or not.

One specific implementation of zero-knowledge proof is called zk-SNARKS (Zero-Knowledge Succinct Non-Interactive Argument of Knowledge). It is a type of proof system where the prover produces a succinct and efficient proof of the correctness and the verifier is fast to verify the proof. The main steps to construct a zk-SNARK are:

- 1) **Arithmetic Circuit:** The function  $f(w, x) = y$  needs to be represented as an arithmetic circuit consisting of multiplication and addition gates. This process is done

<sup>1</sup>Smart-Contracts Dapp:

<https://github.com/pierg/fact-fortress-dapp>

Front-end:

<https://github.com/pierg/fact-fortress-frontend>

Circuit Compiler Tool:

<https://github.com/pierg/fact-fortress-circuits>

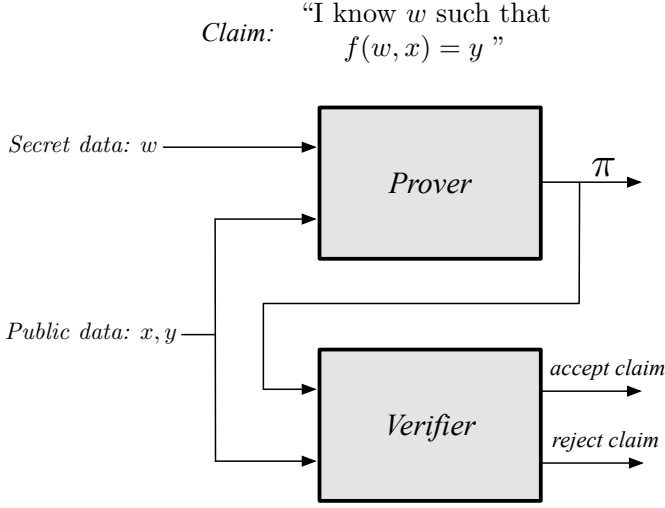


Fig. 1: Generic Prover and Verifier Interaction

by first converting  $f$  into a RICS (Rank-1 Constrain System), which is then transformed into a series of quadratic arithmetic programs (QAPs) using techniques such as Lagrange Interpolation and Fast Fourier Transform (FFT).

- 2) **Setup Procedure:** The setup procedure generates the proving and verifying keys  $S_p$  and  $S_v$ , as well as the public parameters  $pp$ . The prover uses  $S_p$  to generate proofs, the verifier uses  $S_v$  to verify those proofs, and  $pp$  defines the mathematical structure used to construct the circuit and is used by all parties. There are three types of setup procedures:
  - a) *Trusted setup*, where a trusted party generates  $(S_p, S_v, pp)$  for a specific circuit.
  - b) *Trusted but updatable setup*, which is similar to the trusted setup but allows the trusted party to update parameters and keys for other circuits.
  - c) *Transparent setup*, which generates parameters and keys using a publicly known deterministic algorithm, eliminating the need for a trusted party but may be more computationally expensive.
- 3) **Prover:** The prover takes as input the private input  $w$ , the public data  $x$  and  $y$ , the proving key  $S_p$ , and the parameters  $pp$ , and uses them to generate a proof  $\pi$ . This involves constructing a witness polynomial that satisfies the QAP and evaluating it at carefully chosen points, resulting in a proof consisting of two polynomials.
- 4) **Verifier:** The verifier takes as input the public data  $x$  and  $y$ , the proof  $\pi$ , the verification key  $S_v$ , and the public parameters  $pp$ , and uses them to check the validity of the proof. This involves checking that the polynomials in the proof satisfy certain constraints, such as the QAP and the arithmetic relations of the circuit.
- 5) **Zero-Knowledge Property:** Finally, if the proof is valid, the verifier accepts the proof without learning anything about the private input  $w$  or the computation of  $f(w, x) = y$ , except for the fact that the computation is correct.

In summary, the setup procedure generates the proving and verifying keys and the public parameters, while the prover

constructs a proof, and the verifier checks the validity of the proof. If the proof is valid, the zero-knowledge property guarantees that the verifier learns nothing about the private input.

### III. CHALLENGES OF DEPLOYING ZERO-KNOWLEDGE PROOF FRAMEWORKS

Several challenges hinder the adoption and deployment of ZKP frameworks in real-world applications. In this section, we discuss three key challenges: trust, adoption, and scalability.

#### A. Trust

Trust is a critical challenge that must be addressed before zero-knowledge proof (ZKP) frameworks can be widely adopted in the real world. This is particularly important when working with sensitive data, where proving and verifying mechanisms should be transparent and the access to the sensitive data be regulated.

One challenge related to trust is ensuring the authenticity and provenance of the input data used by the prover and the verifier. If the input data used by the prover is compromised or tampered with, the resulting proof may be invalid, compromising the overall security of the system. On the other hand, the verifier needs to be confident that the public data used in the verification process is valid and has not been manipulated to produce a false result.

To address trust-related challenges, one approach is to include the input data as part of the trusted setup procedure. However, this requires trust in the entity that generated the public parameters. This technique can be used to deploy a trusted prover and verifier with fixed trusted data. However, it does not allow for the same prover and verifier to be used on other data, as the setup procedure would need to run again, and a new prover/verifier should be created.

Another approach is multi-party computation (MPC) [2], which allows multiple parties to jointly compute a function without revealing their inputs to each other. We can use MPC to collectively create the input data from several trusted entities. This technique can provide a more robust solution for generating trusted input data and can be used to deploy ZKP systems on a broader range of data without requiring a new trusted setup for each case, however it requires the coordination of multiple trusted entities.

#### B. Adoption

The adoption of ZKP frameworks has been limited, in part, due to the complexity of expressing functions in terms of arithmetic circuits, which are a crucial component of many ZKP protocols. However, over the recent years, several tools and languages have been developed to facilitate the use of ZKP frameworks.

Hardware Description Languages (HDL) such as Circom [3] provide a way to describe circuits in a low-level format that can be compiled to arithmetic circuits. Libraries such as Arkworks [4] provide modular building blocks that can be combined to create custom ZKP systems.

Finally, there are several programming languages designed to make it easier to implement ZKP frameworks in real-world applications. For example, Zokrates [5] is a popular open-source toolkit that allows developers to write programs in a high-level language and compile them into arithmetic circuits. Noir [1] is a Rust-like language that provides tools for ZKP construction and verification. Leo [6] and Cairo [7] are another programming languages that are designed to allow easy expression of complex circuits in a high-level format.

Even with the availability of Domain-Specific Languages (DSLs), Programming Languages (PLs), and libraries, expressing complex functions in terms of arithmetic circuits remains a challenge for non-domain experts.

### C. Scalability

Scalability is a significant challenge in ZKP frameworks, particularly regarding the performance of the prover and verifier with respect to the size of the input data. As input data grows, the time and computational resources required to generate and verify proofs increase significantly, potentially hindering scalability.

TODO, present related: [8], [9], [10], [11]

## IV. FACT FORTRESS

In the following sections we provides an overview of our proposed solution named FACT FORTRESS. In Section V we address the problem of trust, providing general circuit design and smart-contract architecture that has data privacy and authenticity at its core. In Section VI we address the problem of adoption, we propose abstraction layers on top of existing frameworks that facilitate the circuit and data specification. In Section VII we present some scalability results of our approach.

## V. TRUST BY DESIGN

In our approach, the trust of the data is embedded in the design of each circuit as well as in the overall framework. Specifically, each circuit is designed to certify the proof of provenance of the data, ensuring that the input data is coming from authenticated sources. Moreover, the overall architecture framework is designed to facilitate the exchange of data among different parties in a regulated way implementing access control policies.

By employing this our approach, the trustworthiness of the ZKP framework is enhanced, making it more suitable for handling sensitive data and real-world applications.

### A. Proof of Provenance

Our approach involves embedding a ‘*proof of provenance*’ alongside the proof of statement in each circuit. Figure 2 provides an example of how we prove the truthfulness of a generic statement ‘*I know  $w$  from  $A$  such that  $f(w) = y$* ’, without revealing any information about  $w$ . The proof is constructed in two steps:

Step 1: Proof of provenance, which proves that:

- $w$  originated from the authority  $A$

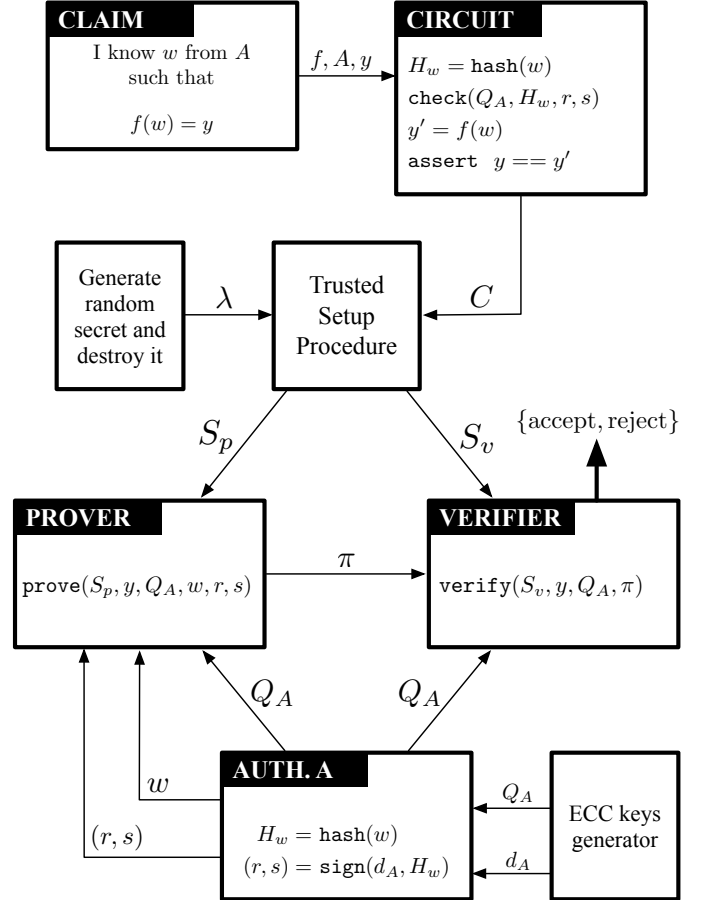


Fig. 2: Statement Verification Process using ZKP

- $w$  has not been tampered with or altered in any way

Step 2: Proof of statement, which proves that:

- The function  $f$  has been faithfully translated into an arithmetic circuit
- The result of  $f(w)$  is equal to the claimed result  $y$

1) Protocol: TODO

### B. Smart-Contracts Architecture

To ensure transparency and accountability in data access, our FACT FORTRESS framework incorporates smart contracts into its architecture. The overall architecture of the framework is depicted in Figure 3. By incorporating smart contracts into the architecture, we ensure that all parties involved in the data analysis adhere to the specified policies and that the public results of the analysis can be trusted.

The framework allows certified *data providers* to securely store their sensitive data and set *data access policies* on how the data must be handled. Data analysts can request access to the data based on these policies to perform an analysis and compute the ZKP locally. Alternatively, they can delegate the data analysis to FACT FORTRESS, which returns the zero-knowledge proof of the computation and the result directly to them.

The types of analyses that can be performed are defined by a library of functions that can be computed on data of

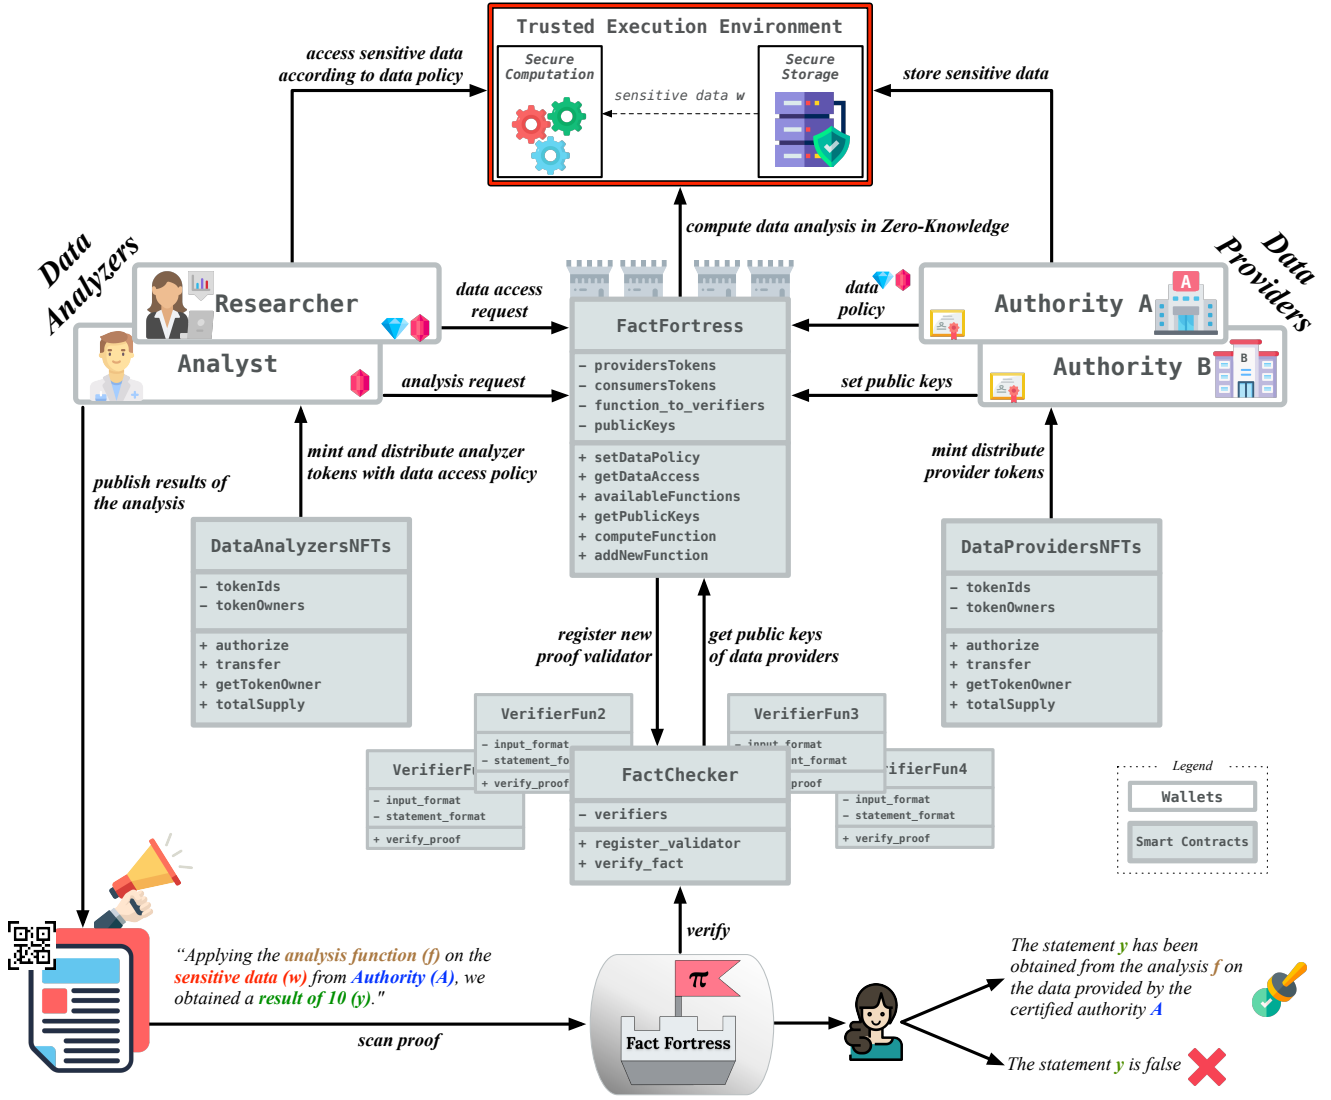


Fig. 3: Smart Contracts Architecture on Blockchain

any form. For each function, we deploy a verifier on-chain, which anybody can use to validate a proof. When a proof is submitted for validation, FACT FORTRESS dispatches it to the correct verifier, which checks the following:

- 1) The proof was generated by the function claimed by the analyst.
- 2) The data used to generate the proof has not been tampered by the analyst and it comes from a certified data provider.
- 3) The claimed result is the correct result of the function applied to the data.

By using smart contracts, we can guarantee the integrity and transparency of the data access policies and their enforcement, which is essential in sensitive domains such as healthcare or finance.

1) *Data-Access Policies:* To provide secure and regulated data access, our framework incorporates data-access policies. Each data provider can define their access policies, which include the type of data that can be accessed, the duration of access, and any restrictions on data usage.

Data analysts can request access to the data by providing

a proof of compliance with the access policies. If the proof is valid, they receive a non-fungible token (NFT) that grants them temporary access to the requested data.

Our access control mechanism ensures that only authorized data analysts can access the data and that they can only use the data according to the specified policies. The NFTs also ensure that data analysts cannot access data that they are not authorized to use or extend their access beyond the specified duration.

Our framework leverages the use of a transparent and publicly verifiable ledger, such as a blockchain, to ensure that data usage is logged and monitored in a transparent and immutable way. This creates an audit trail that tracks all data access and usage on the ledger, allowing data providers to monitor how their data is being used and to detect any unauthorized access or usage. By ensuring the integrity and transparency of data access and usage, our framework provides a robust and trustworthy solution for data sharing and analysis.

2) *Protocol:* TODO

3) *Ethereum Implementation:*



## VI. ABSTRACTING CIRCUITS

To increase adoption of zero-knowledge proof (ZKP) frameworks and enable developers with limited experience to implement functions in real-world applications, more user-friendly interfaces and higher-level abstractions are needed to abstract away the low-level details of arithmetic circuits.

One approach to achieve this is through the development of tools for automatic circuit synthesis. We have implemented such a tool by building abstraction layers on top of the Noir framework [1]. Our abstractions provide simple Python APIs that allow developers to express complex computations more easily, without requiring them to understand the underlying arithmetic circuits.

Our tool (code available<sup>2</sup>) enables the automatic generation of arithmetic circuits from high-level abstractions, making it easier to implement ZKP protocols in their applications.

Our library provides clear and abstract APIs that allow users to specify the data format, the function to be performed by the circuit, and the authority that provided the data. The library compiles down from Python API to JSON configuration file and ultimately parses the JSON and compiles the fully functioning circuit in Noir as shown in Figure 4.

In order to produce a circuit our library takes as input:

- *data*: arrays of any size and shape, elements can be integers, string or double. Doubles are quantized by our library with the specified specision.
- *authority*: private keys of the authority providing the data
- *function*: analysis that must be performed on the data. This function must be chosen by our library, e.g. one can do *average of dot products*, *weighted sums* etc.. and compose multiple primitive functions together.

Our library also supports the automatic generation of data of the desired shape and the creation of fake authorities for demo purposes.

Our library performs the following operations:

- 1) Compute the hash of data, sign the hash using an authority's private key using Schnorr signature protocol.
- 2) Perform the chosen function on the data in Python to compare the result with the one executed by the circuit in Zero Knowledge.
- 3) Generate a new comprehensive configuration file in JSON format to programmatically share and re-create the circuits.
- 4) Generate the circuit! Given a configuration file, our library will generate a structured folder with all files needed to generate the proof in Zero-Knowledge using noir. The circuit compiles and generates valid proofs right out the box without having the user writing anything in any domain specific language. Specifically, our library can generat circuits on any data size and shape and can prove:
  - Proof of Provenance: Compiles circuits that can compute the data hash using SHA256 and checks that the

hash is valid and that the data comes from the authority using Schnorr Signature.

- Proof of Statement: Translates the chosen function in a valid circuit and checks that the function applied on the data results in the expected statement previously computed in Python.

Once the process has completed, the user can navigate to the generated folder and run the following commands to prove and verify the circuit respectively:

```
# PROVE
nargo prove p

# VERIFY
nargo verify p
```

## VII. ADDRESSING SCALABILITY

Scalability is a critical challenge for any ZKP framework as the proving and verification times typically increase with the size of the input data and the number of arithmetic gates required by the function.

To address this challenge, we have conducted experiments on one of the function from our library named `average_dot_products` having as input a two-dimensional matrix and a vector. In this function, the dot-product between each row of the matrix and the vector is computed and then averaged over all rows.

We generated data for matrices and vectors of different sizes, where each element is a random integer of 16 bits. Figure 5 shows the proving and verification times on different sizes of data. Figure 6 shows the growing number of operations as the data size increases, where the growth is linear.

TODO...

## REFERENCES

- [1] "Noir," <https://noir-lang.org>.
- [2] O. Goldreich, "Secure multi-party computation," *Manuscript. Preliminary version*, vol. 78, no. 110, 1998.
- [3] "Circom," <https://docs.circom.io>.
- [4] "Arkworks," <https://arkworks.rs>.
- [5] "Zokrates," <https://zokrates.github.io>.
- [6] "Leo," <https://leo-lang.org>.
- [7] "Cairo," <https://cairo-lang.org>.
- [8] H. Wu, W. Zheng, A. Chiesa, R. A. Popa, and I. Stoica, "DIZK: A distributed zero knowledge proof system," *Proceedings of the 27th USENIX Security Symposium*, pp. 675–692, 2018.
- [9] E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev, "Scalable Zero Knowledge with No Trusted Setup," *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, vol. 11694 LNCS, pp. 701–732, 2019.
- [10] E. Ben-Sasson, A. Chiesa, E. Tromer, and M. Virza, "Scalable Zero Knowledge Via Cycles of Elliptic Curves," *Algorithmica*, vol. 79, no. 4, pp. 1102–1160, 2017.
- [11] C. Weng, K. Yang, J. Katz, and X. Wang, "Wolverine: Fast, scalable, and communication-efficient zero-knowledge proofs for boolean and arithmetic circuits," *Proceedings - IEEE Symposium on Security and Privacy*, vol. 2021-May, pp. 1074–1091, 2021.

<sup>2</sup>Configuration file:

[https://github.com/pierg/fact-fortress-circuits/blob/main/circuits/average\\_dot\\_products/config.json](https://github.com/pierg/fact-fortress-circuits/blob/main/circuits/average_dot_products/config.json)

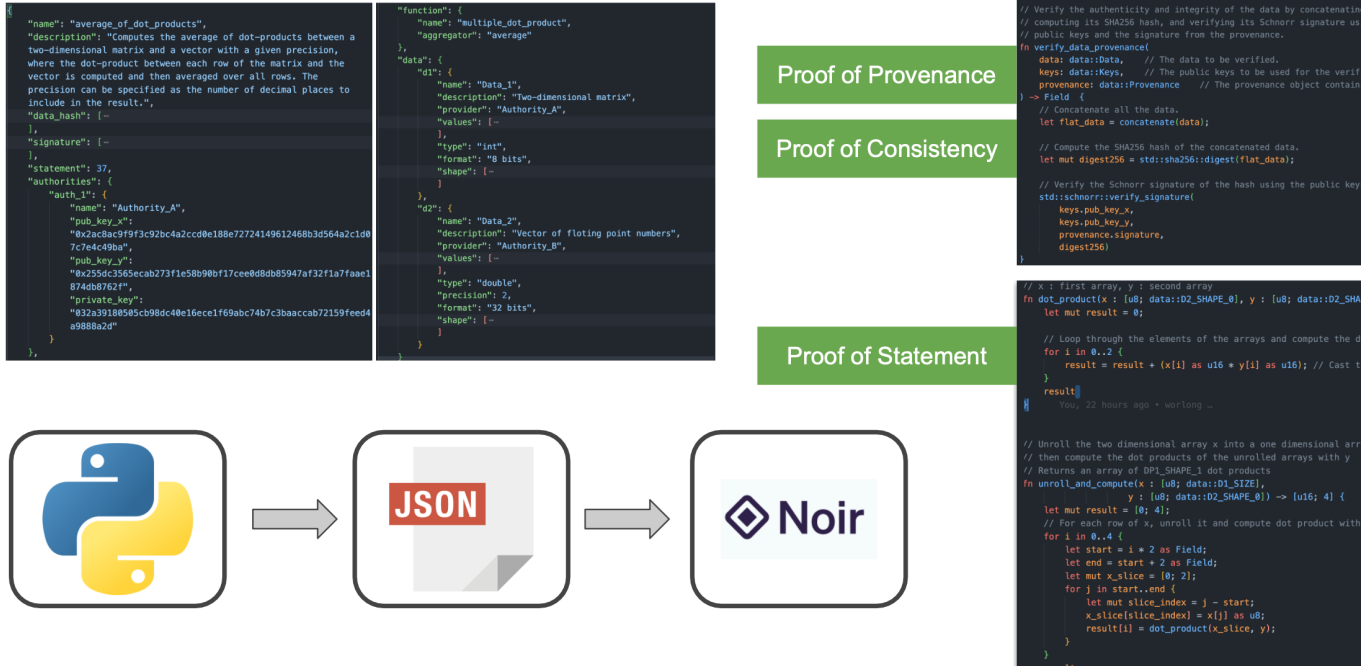


Fig. 4: Circuits Generation Process




Fig. 5: Scalability of prover and verifier




Fig. 6